



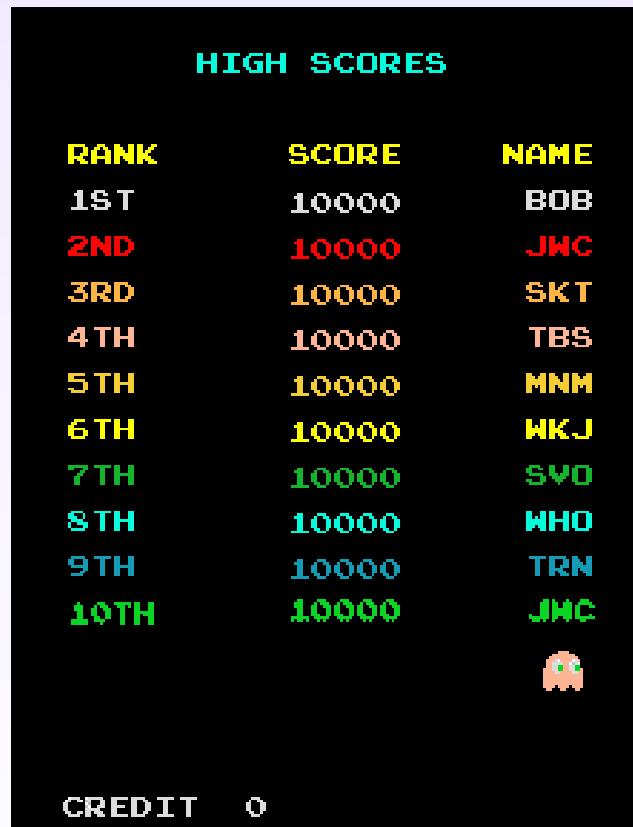
# About Me



PANACHE

- Wiebe-Marten Wijnja
- [@WiebeMarten](#) / [Qqwy](#)
- Creative Developer, Open Sourcerer, Dreamer
- [elixirforum.com](#) Moderator
- [panache-it.com](#)
- Consulting & Placement
- Ruby, PHP, Elixir, ...
- “Panache” *Bold, In a league of its own*

# A Winner is You: A Practical introduction to Concurrency in Elixir



# What this talk is about

- What type of language is Elixir
- Why is concurrency important?
- How is concurrency different in Elixir?
- **Practical Example:**
- How do I make a concurrent application in Elixir?
- How do I make it fault-tolerant, scalable?

Assumed knowledge: Basic understanding of programming.

# Tick Tock...



- Lot of ground to cover for 45 minutes.
- Both general overview and practical example.
- **Not** going into all the syntax.
  - (The docs are amazing!)
- Explanation + Demonstration

# What kind of language is elixir ?

- **Concurrent** ← *Topic of this talk*
- Dynamically typed
- Immutable + Functional
- Built on top of Erlang: 29+ years of Robustness
- Flexibility, Macros
- Explicit over Implicit
  - “Instead of Magic (cool stuff happens but you don't know how or why), Elixir gives you Power (cool stuff happens, and you know **exactly** how and why)”

# Why is concurrency important?

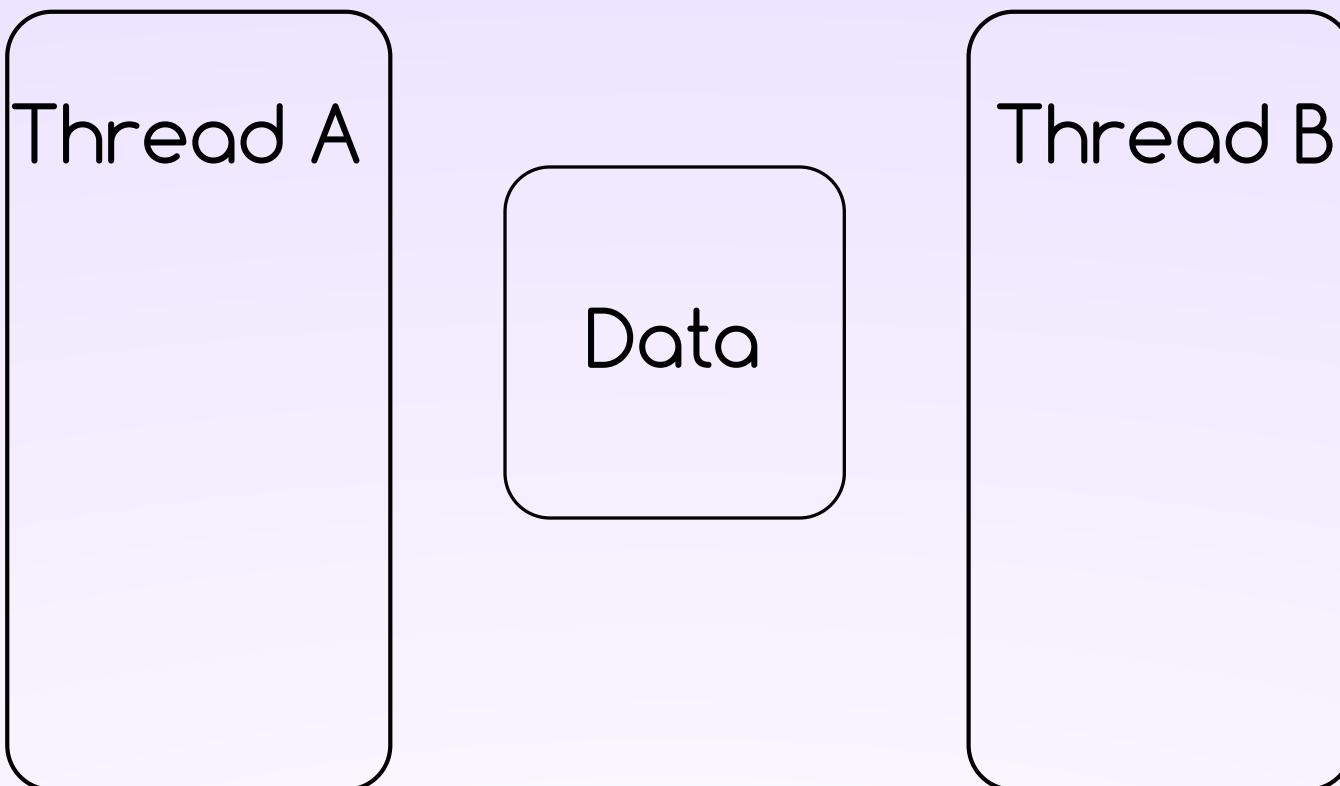
**Concurrent:** “Happening at the same time”

- **Speed**
  - Multicore Computers
  - While waiting for something, do something else as well

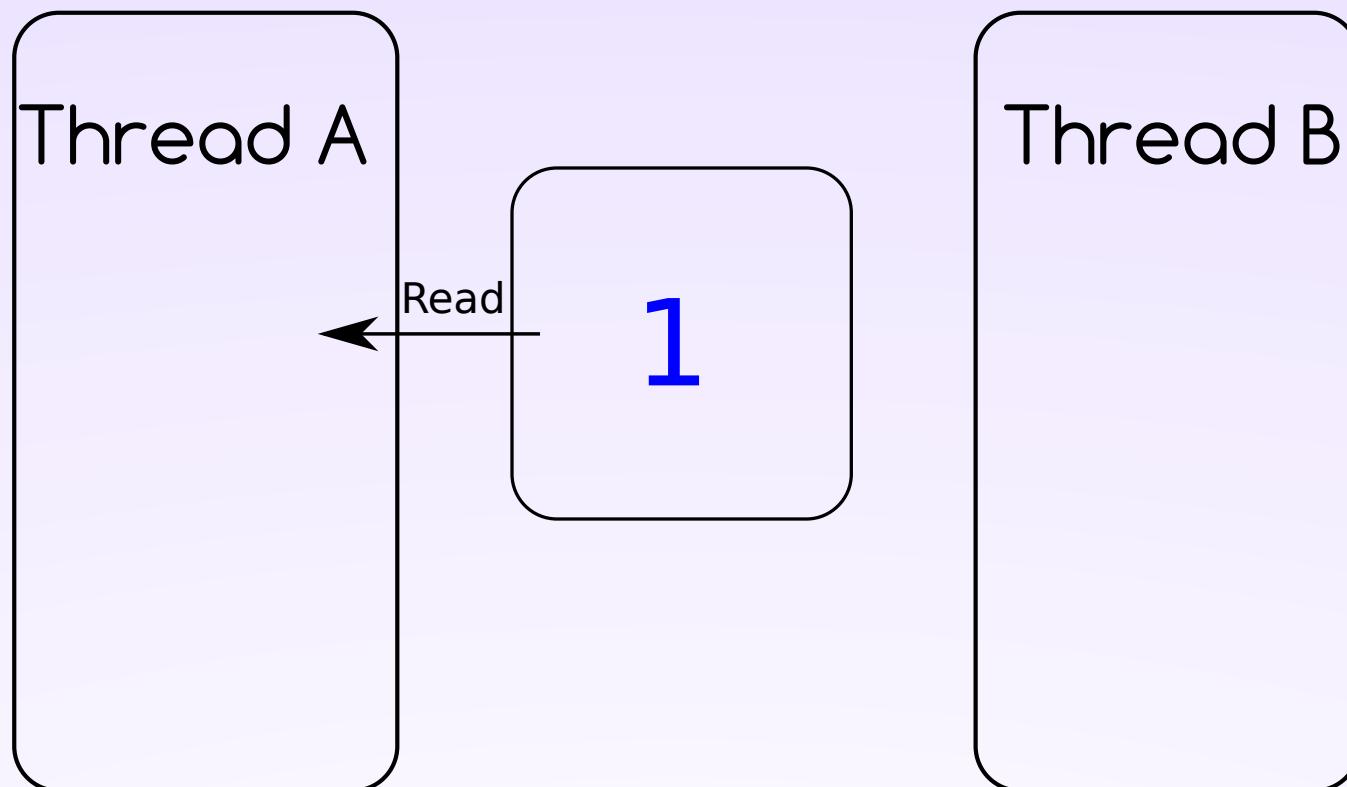
# Why is concurrency hard?

- Multiple things happen at the same time
- Execution order cannot be known
  - Hard to reason about

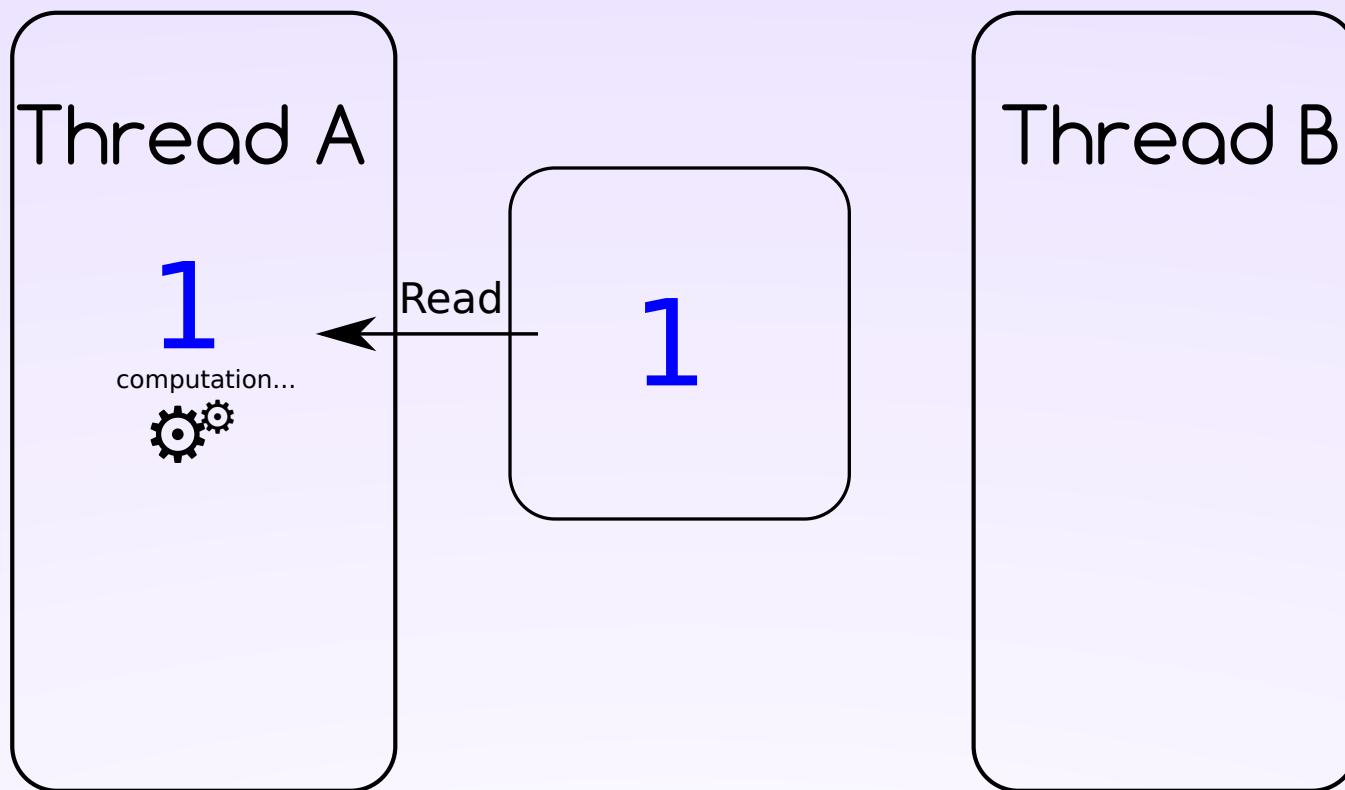
# Why is concurrency hard?



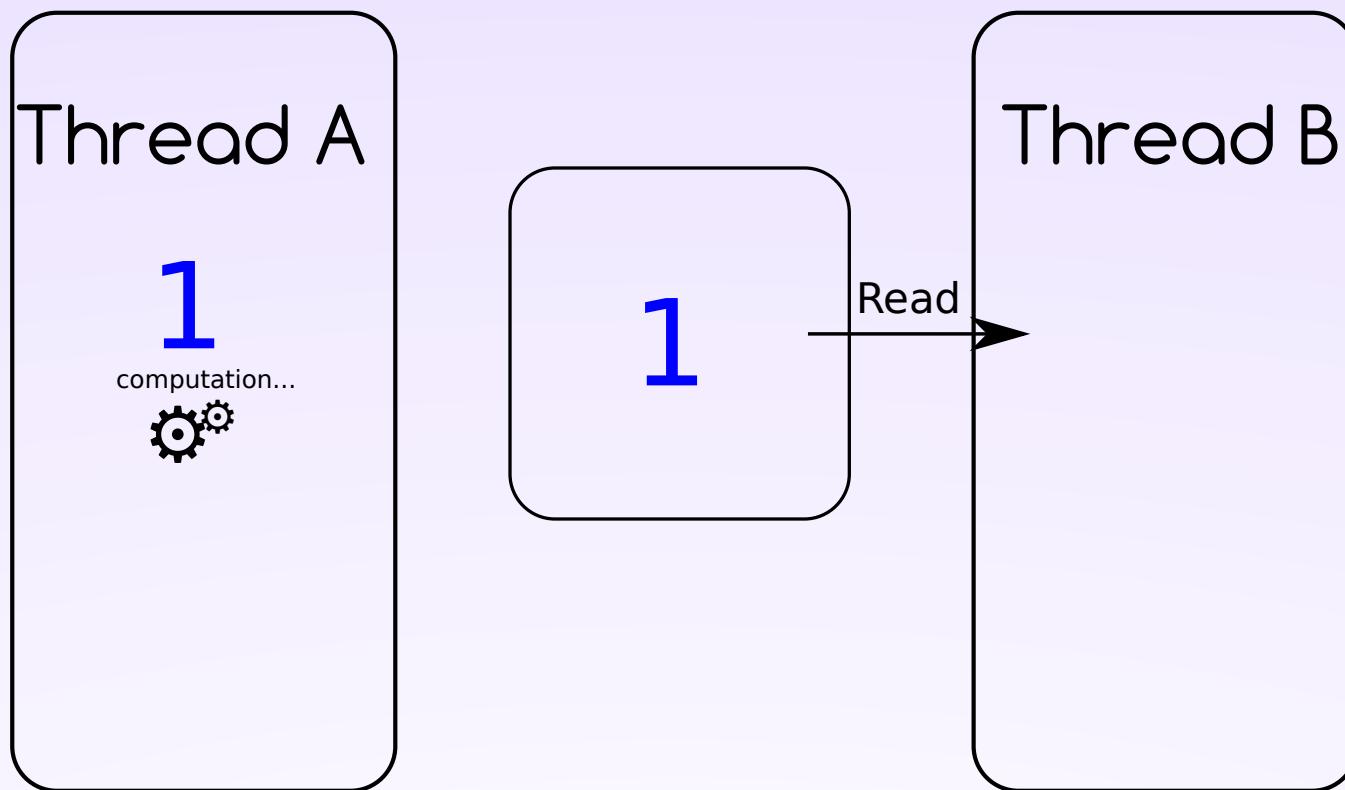
# Why is concurrency hard?



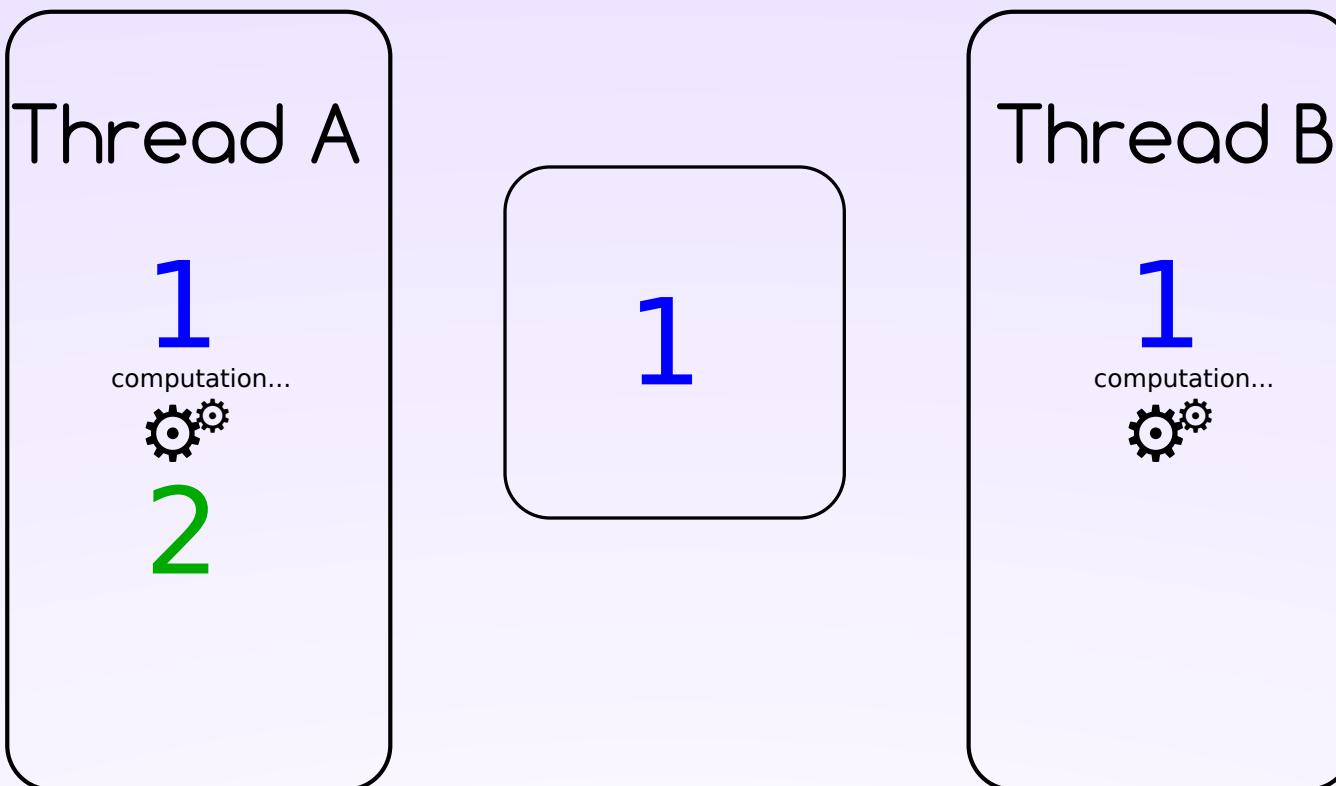
# Why is concurrency hard?



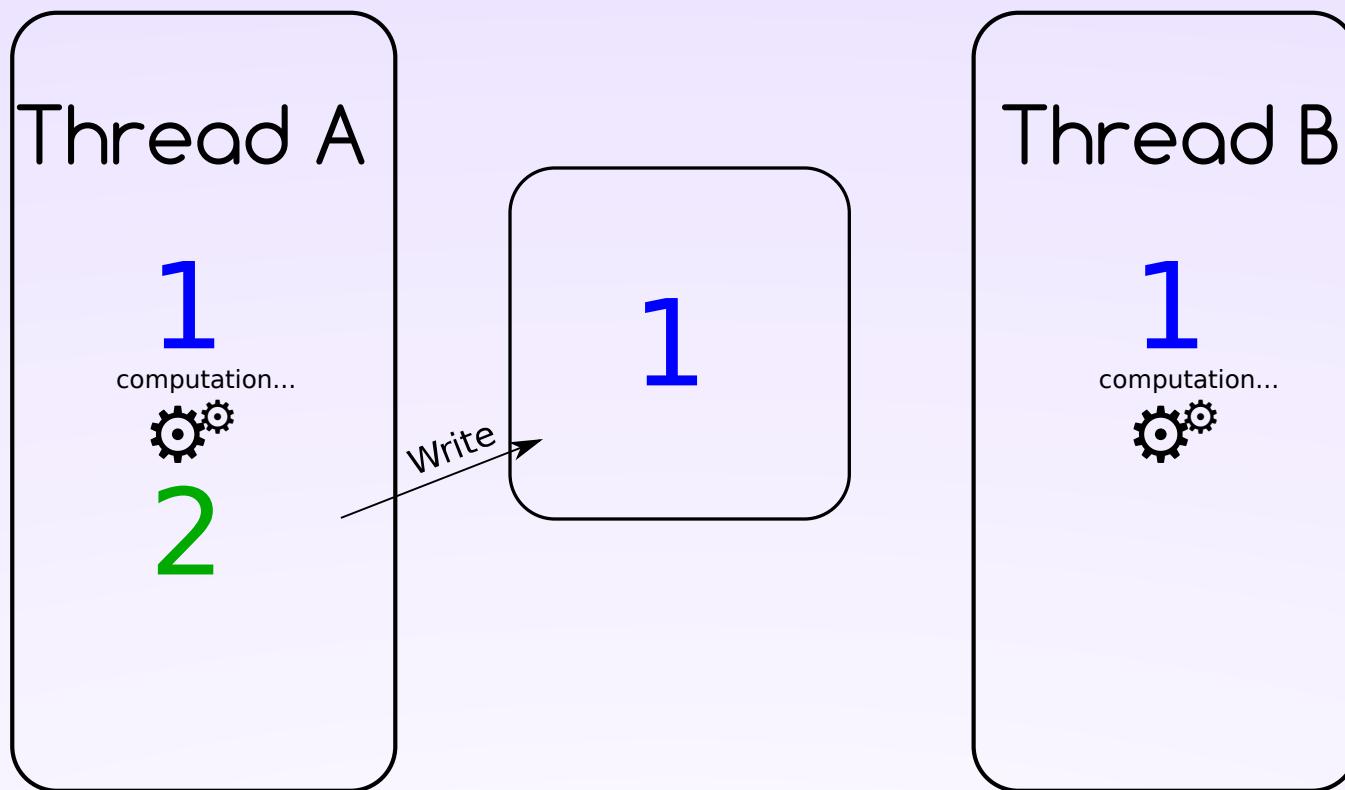
# Why is concurrency hard?



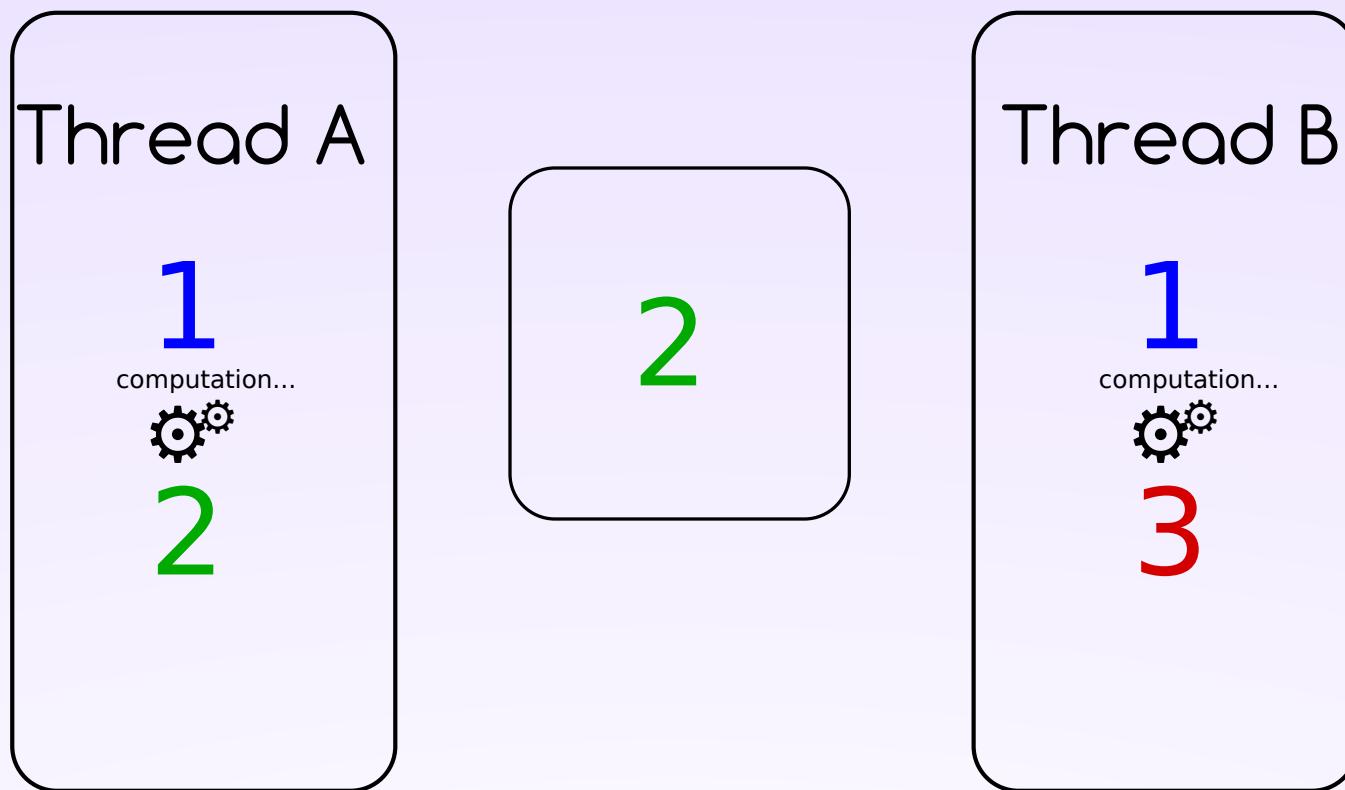
# Why is concurrency hard?



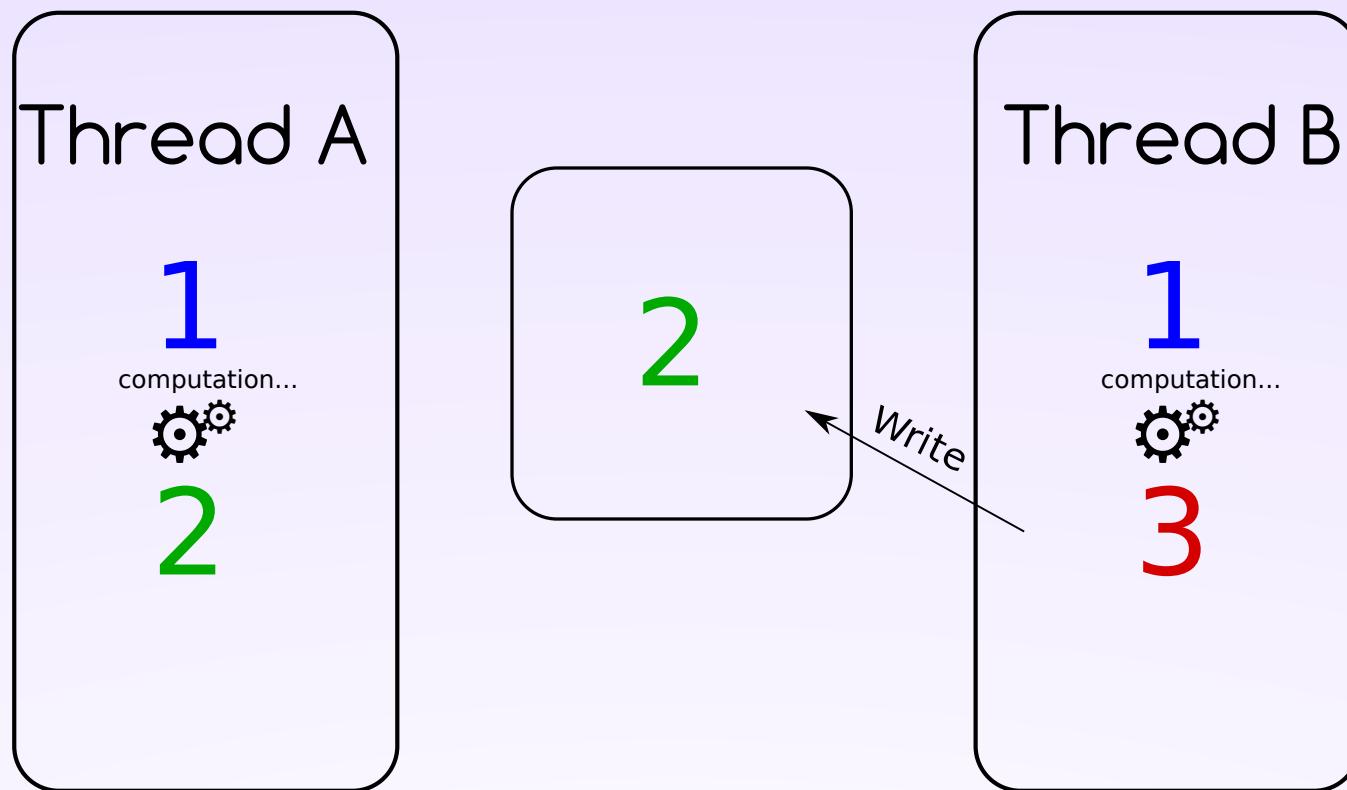
# Why is concurrency hard?



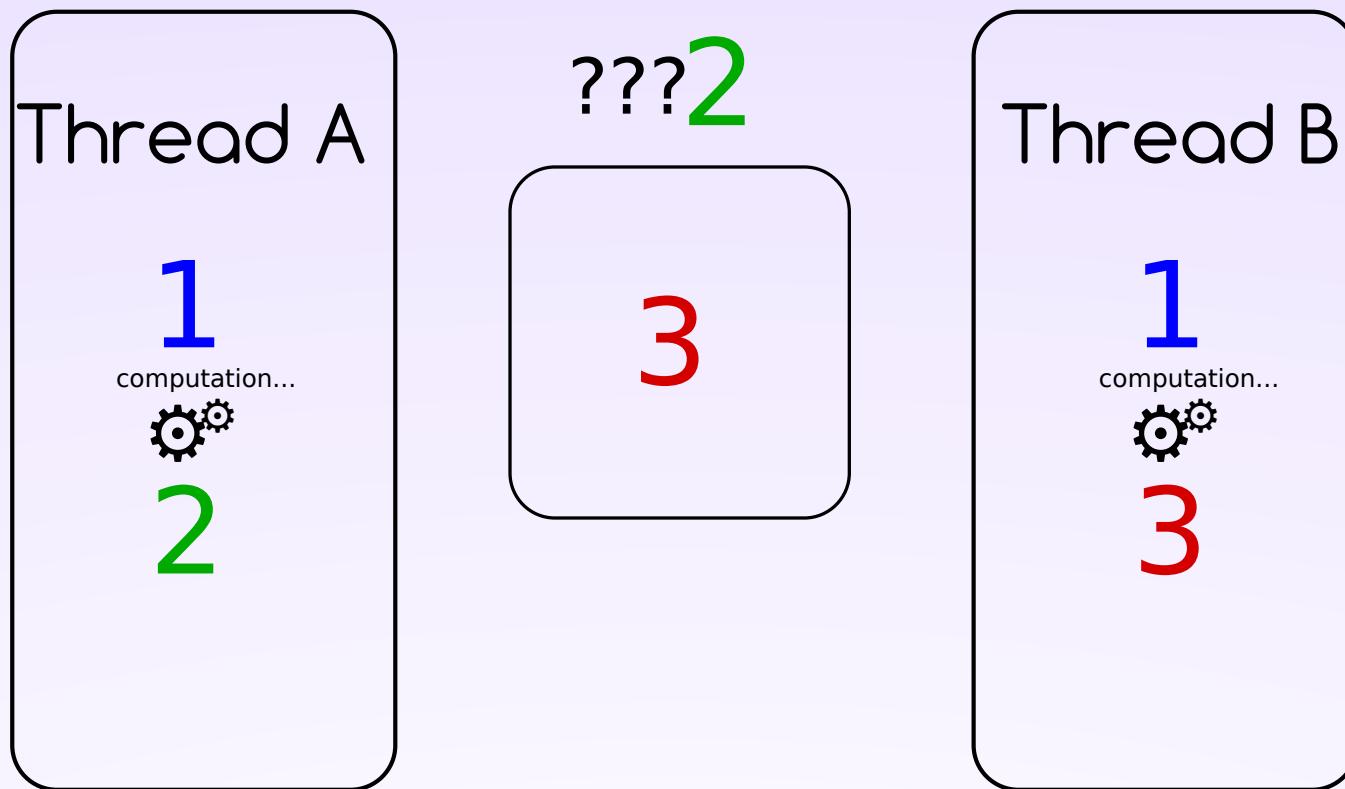
# Why is concurrency hard?



# Why is concurrency hard?

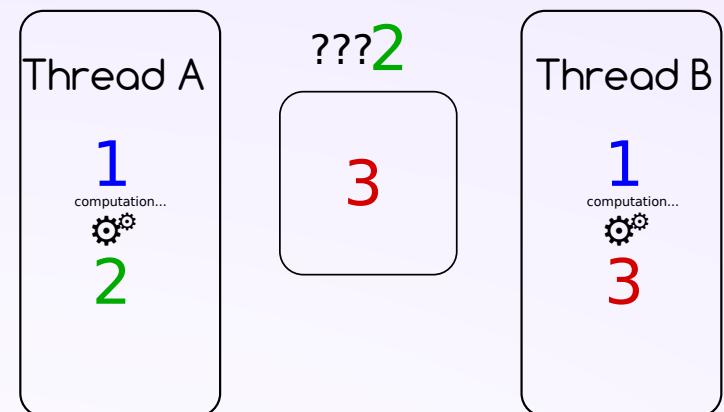


# Why is concurrency hard?



# Why is concurrency hard?

- Race Conditions
- Locks, Mutexes, etc. are only band-aids
- Very hard to reason about



**The only real solution:**

**Don't share memory between threads!**

# The Actor-Model

- 1) A Process is a lightweight VM-thread with **isolated memory**.
- 2) The only way for processes to exchange information, is by **message passing**.
- 3) A process can spawn new processes, whenever necessary.

This is enforced *by the virtual machine!*



Memory Safety! Speed! No Stop-The-World Garbage Collection!

# The Actor-Model

- **Concurrency:** Possible to *truly* run code at the same time.
- **Distribution:** Possible to distribute to other computers. (message-passing works exactly the same)
- **Fault-Tolerance:** When one part breaks, the rest of the system is unaffected.

# The Actor-Model

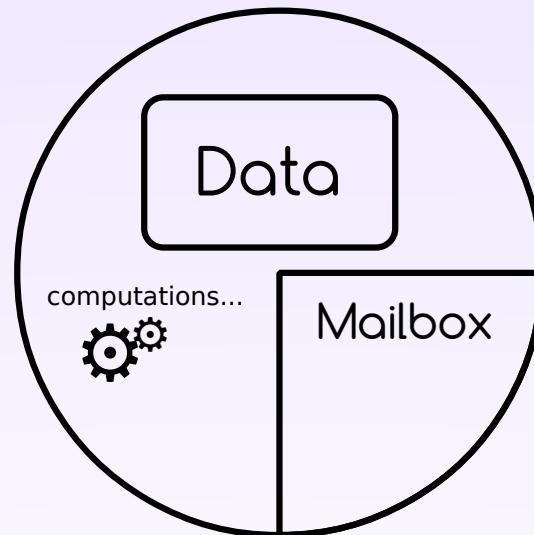
- Programming for Ant Colonies



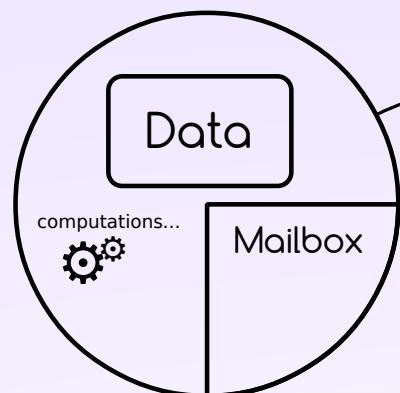
©Warren Photographic

# Meet the Process

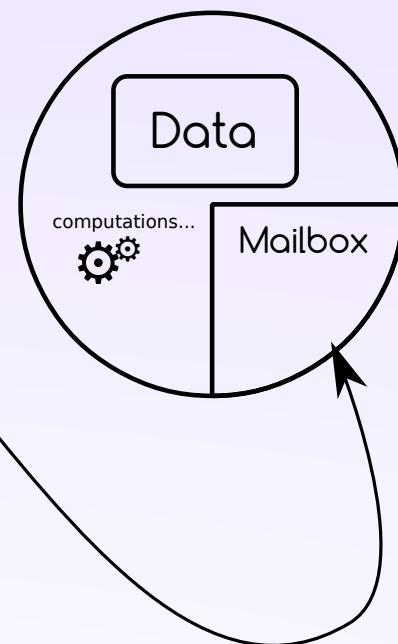
## A Process



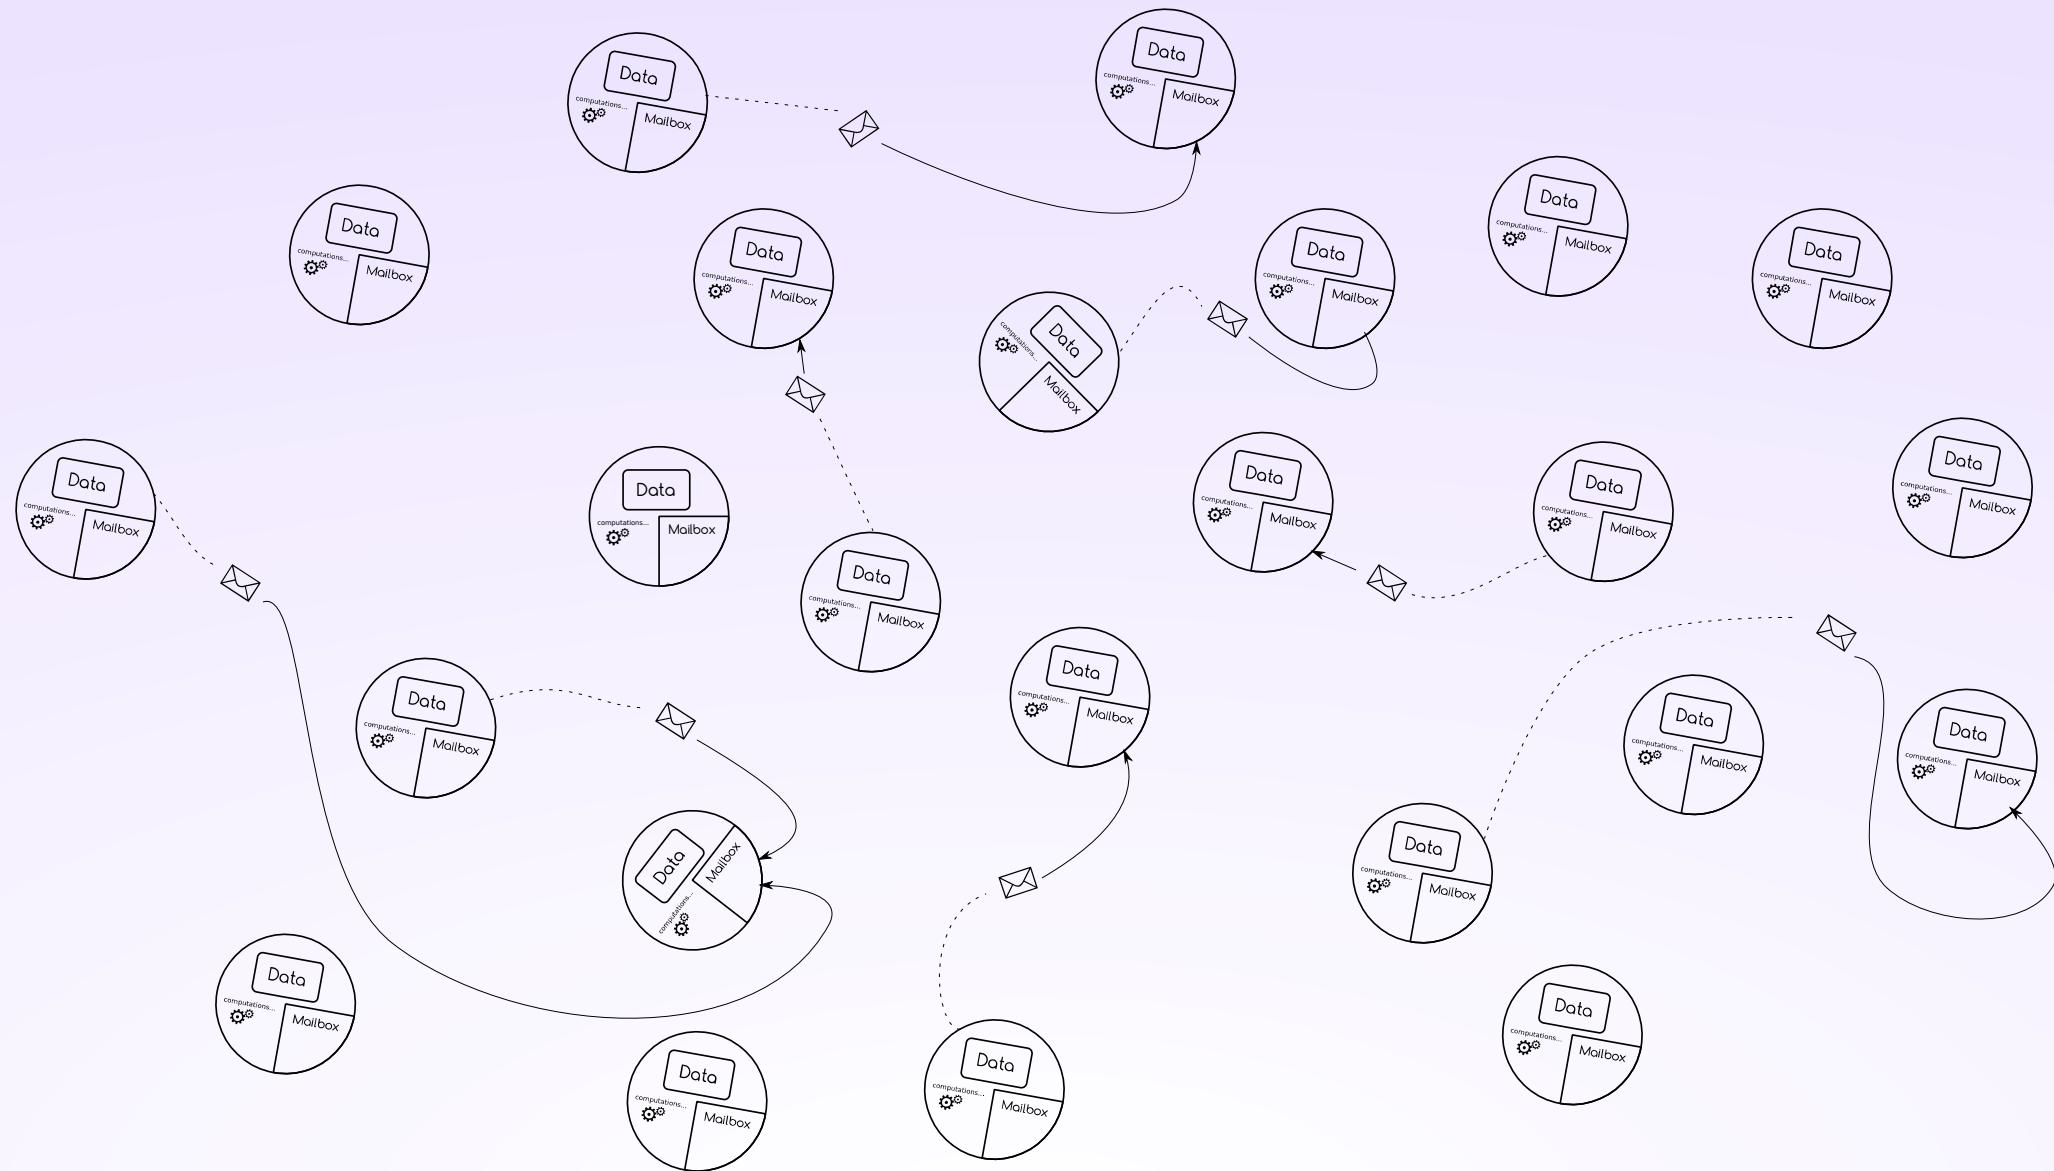
# Process A



# Process B

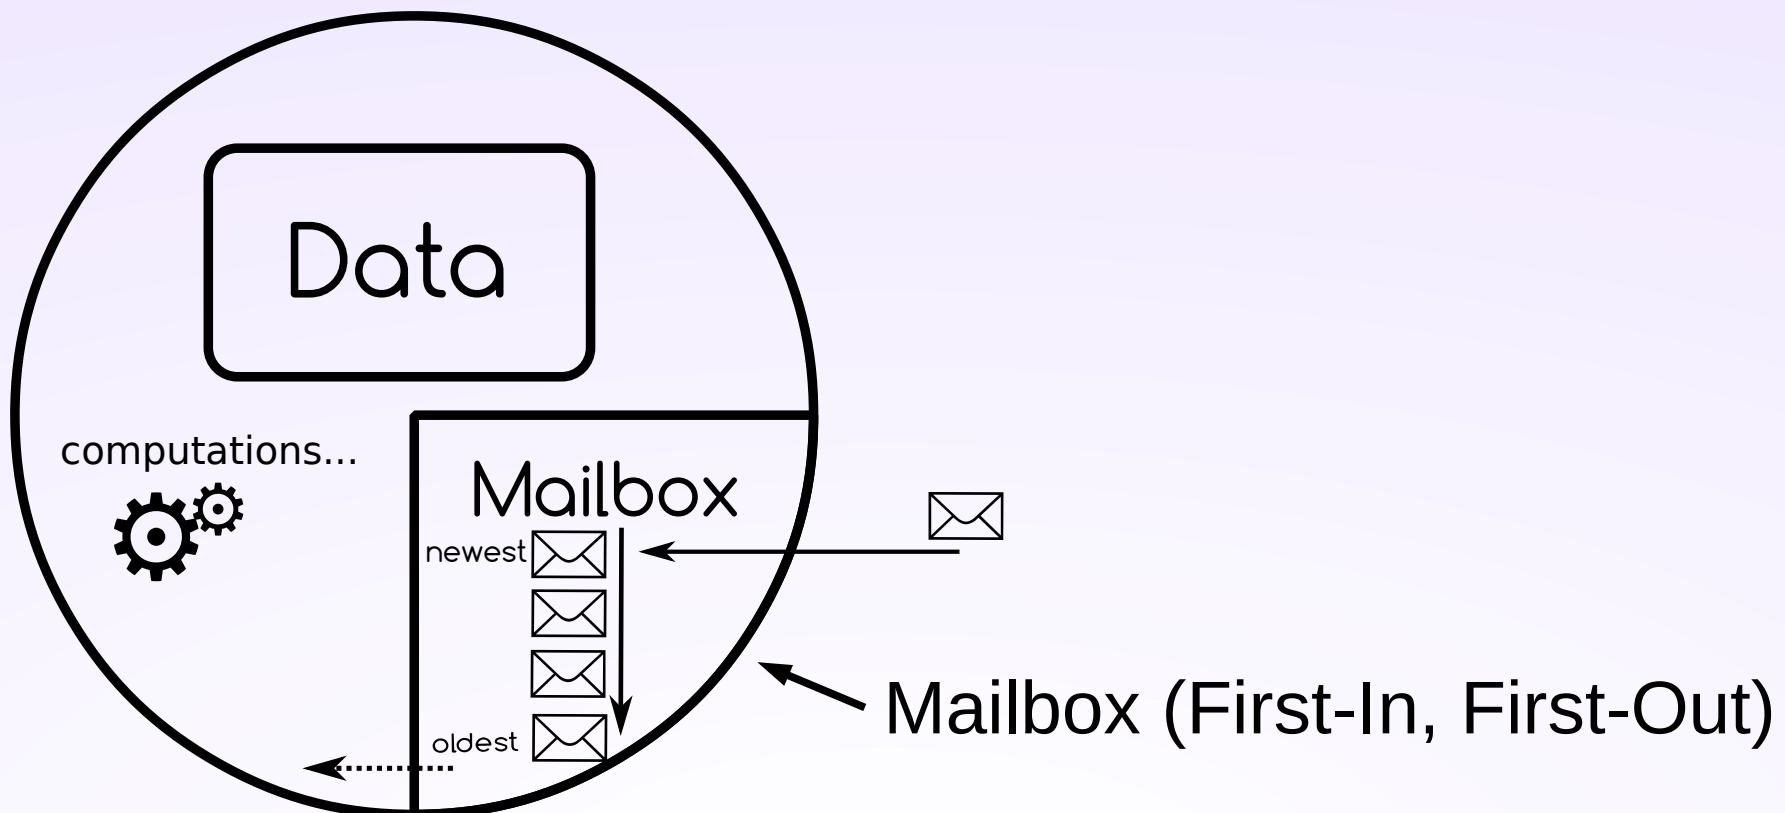


# An Elixir Application:

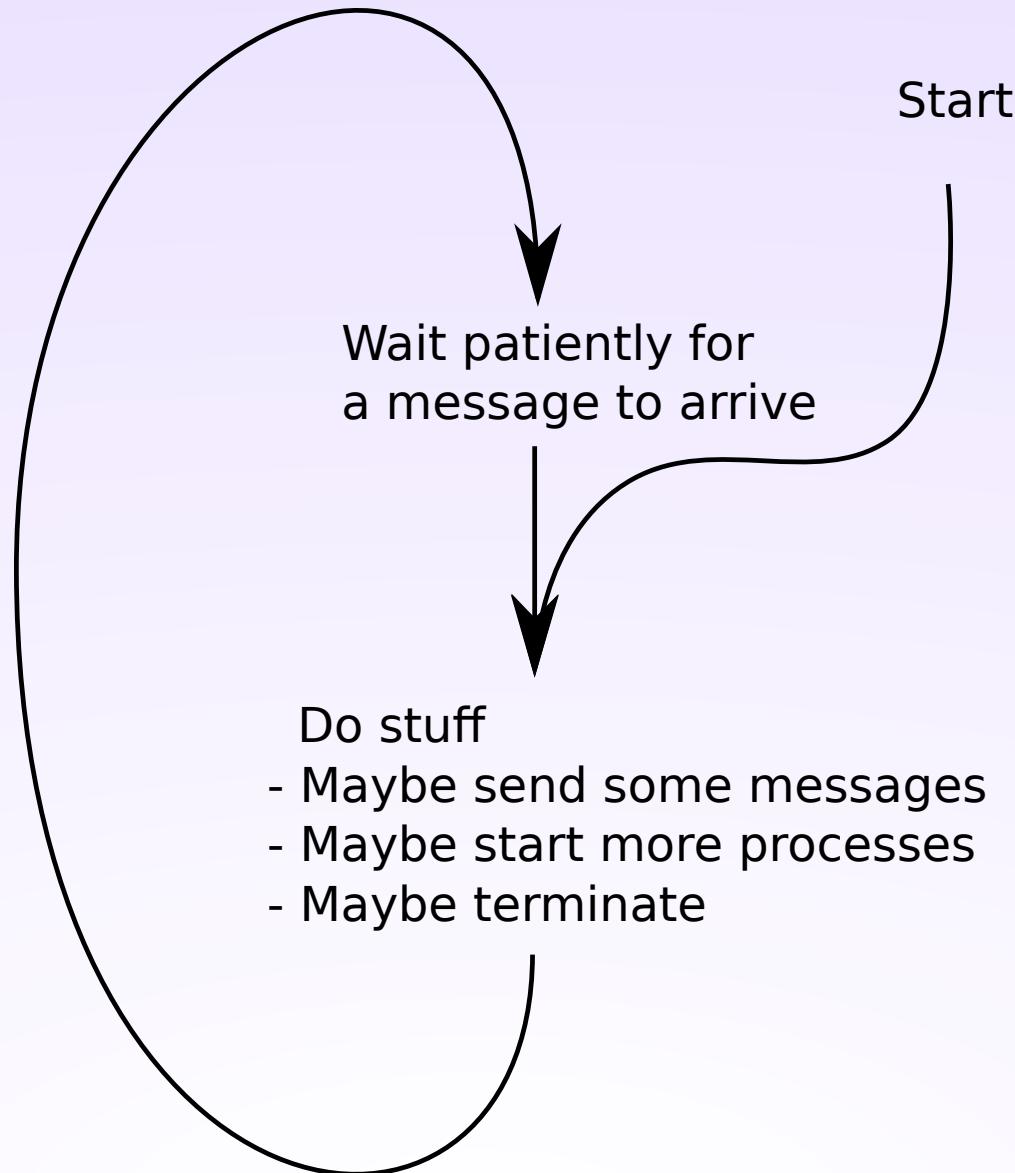


# Anatomy of a Process

<0.124.0> ← Process ID (PID)



# Life of a Process



# Basic building blocks

- `self` → Get the current Process' PID.
- `spawn` → Spawns a new process that executes the given code.
- `send(pid, message)` → Sends a message to a process.
- `receive do` → Waits for the next message (and optionally continues after a timeout)

# Very Basic Example

```
other_pid = spawn(fn ->
  receive do
    {sender, contents} ->
      send(sender, "You sent me: #{inspect(contents)}")
  end
end)

send(other_pid, {self, "Hello, world!"})

receive do
  message -> IO.puts "Main Process received: #{inspect(message)}"
end
```

# OTP

- “Open Telecom Platform”

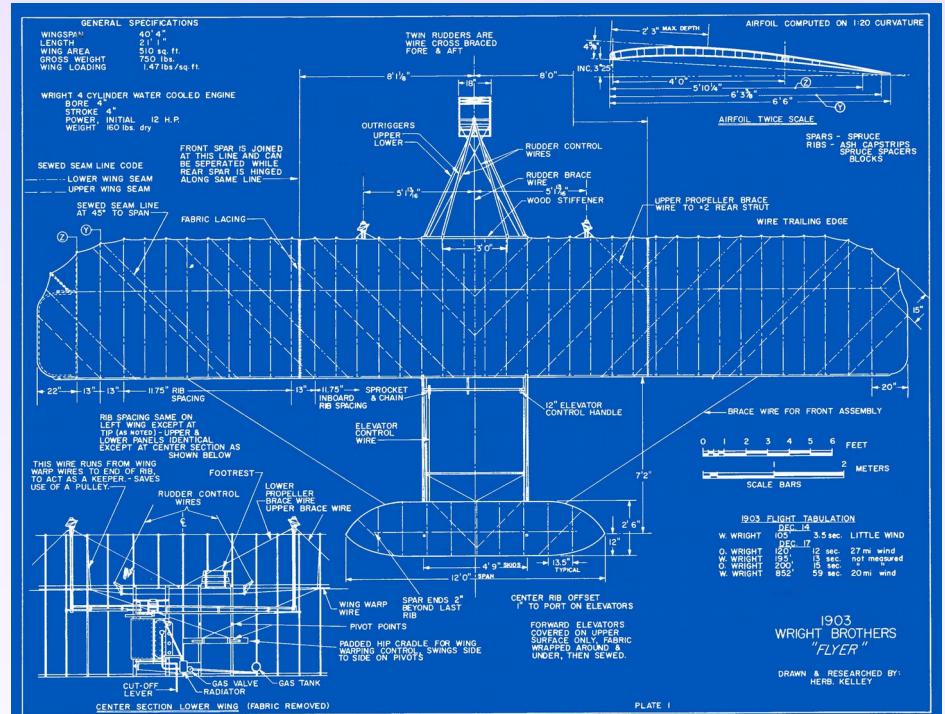
Telecommunication: *“The science and technology of the communication of messages over a distance using electric, electronic or electromagnetic impulses.”*

- Tools to make working with multiple processes easier:
  - Less boilerplate.
  - **Generic** blueprints for often-used kinds of processes.
  - Adds **synchronous** message passing.
  - **Supervision**
  - **Name registration** for processes.
  - **Debugging**
  - *More advanced features:*
    - Hot code reloading
    - Distributed failover/takeover

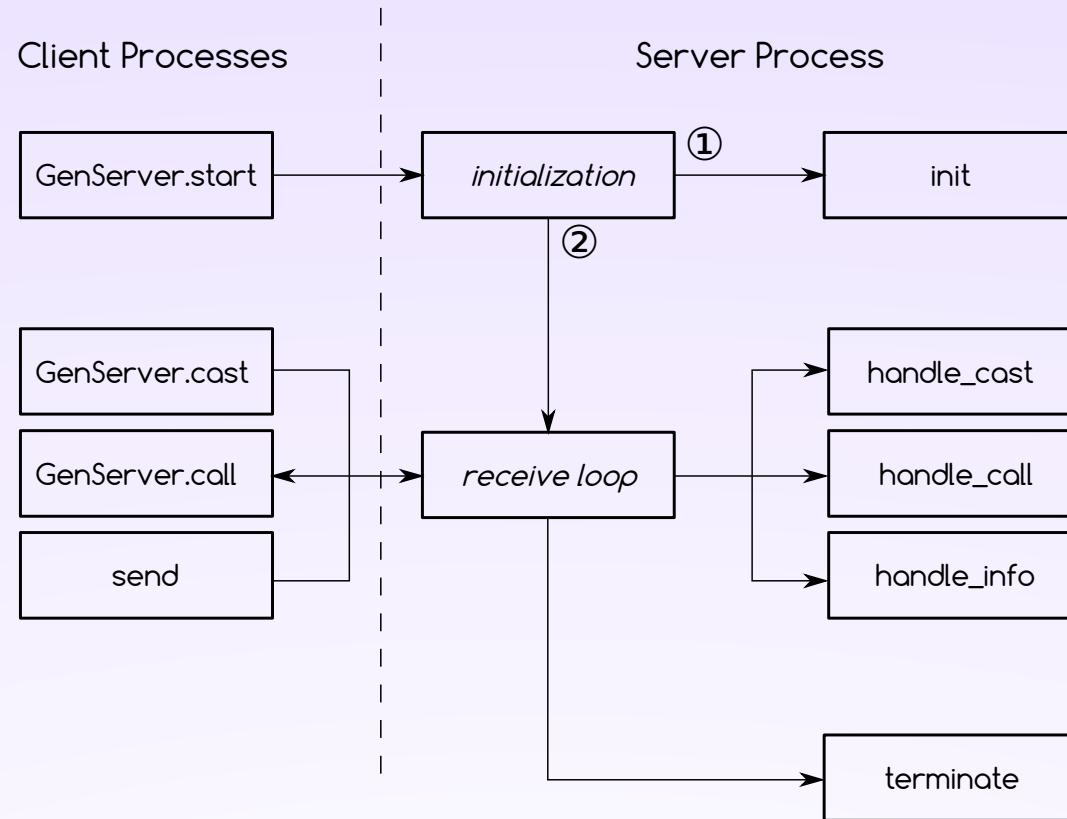


# Generic Blueprints

- **Agent:** Simple wrapper around state.
- **Task:** Single-use concurrent computations.
- **GenServer:** Basic process behaviour blueprint that supports:
  - Initialization
  - Both sync and async message passing
  - Cleanup
  - Both *Agent* and *Task* are built on top of this.
- Some others that are used less frequently (GenEvent, gen\_fsm)

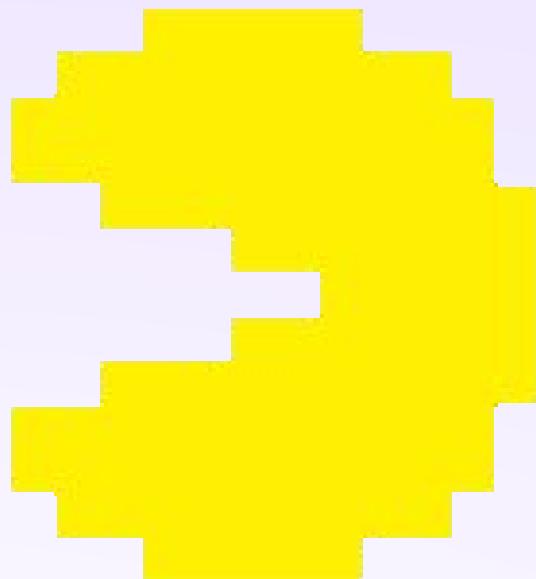
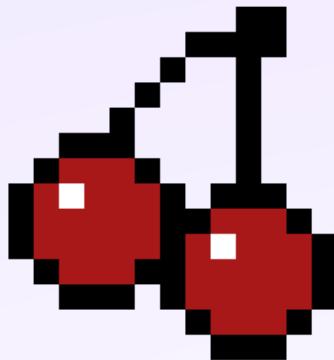


# GenServer Life Cycle



# Practical Example: High Scores Table

- Easy to understand
- Not too trivial
- Not a TODO-list



# Practical Example: High Scores Table

- Save the top X player's scores.
- For each score, save:
  - player name
  - score
  - time
- Return ordered top X when requested.

HIGH SCORES		
RANK	SCORE	NAME
1ST	10000	BOB
2ND	10000	JHC
3RD	10000	SKT
4TH	10000	TBS
5TH	10000	MNM
6TH	10000	WKJ
7TH	10000	SVO
8TH	10000	WHO
9TH	10000	TRN
10TH	10000	JHC

 CREDIT 0

# HighScoresTable GenServer

- Internal state:
  - list of scores (ordered),
  - Max. list length
- `init` → Store how many players we should keep track of.
- `add_score` → adds score to list if higher than lowest.
  - Returns `{:congrats, position}` if the player made it in the high scores
  - Returns `:unfortunate` if the player did not.
- `current_scores` → Returns the current list of scores.

*-Live Demonstration-*

```

defmodule HighScoresTable do
  use GenServer
  @moduledoc "A process that stores a table of high scores."
  @doc "A table is list of scores and its max. length"
  defstruct scores: [], max_players: nil

  defmodule Score do
    @doc "A single score"
    defstruct name: "", score: 0, time: nil
  end

  # Outward API

  @doc "Starts the process, internally calls `init`."
  def start_link(max_players) do
    GenServer.start_link(__MODULE__, max_players)
  end

  @doc "Adds a new `score` for `name` to the table."
  def add_score(pid, name, score) do
    time = :erlang.time |> Time.from_erl!
    score = %Score{name: name, score: score, time: time}
    GenServer.call(pid, {:add_score, score})
  end

  @doc "Lists current scores"
  def current_scores(pid) do
    GenServer.call(pid, :current_scores)
  end

  # internal GenServer callbacks

  @doc "Initializes the GenServer"
  def init(max_players) do
    initial_state = %HighScoresTable{max_players: max_players}
    {:ok, initial_state}
  end

  @doc "Adds a score to the table, updating the internal state."
  def handle_call({:add_score, score = %Score{}}, _from, state) do
    {response, new_scores} =
      maybe_add_score(state.scores, state.max_players, score)
    new_state = %HighScoresTable{state | scores: new_scores}
    {:reply, response, new_state}
  end

  @doc "Lists the current scores"
  def handle_call(:current_scores, _from, state) do
    {:reply, state.scores, state}
  end

  defp maybe_add_score(scores, max_players, new_score) do
    split_scores = Enum.split_while(scores, fn score ->
      score.score >= new_score.score end)
    case split_scores do
      {all, []} when length(all) >= max_players ->
        {:unfortunate, scores}
      {higher, lower} ->
        new_scores = (higher ++ [new_score] ++ lower) |>
          Enum.take(max_players)
        {{:congrats, length(higher)+1}, new_scores}
    end
  end
end

```

# It works!

But...

# Hmm...



How do we make  
this:

- Fault-tolerant?
- Scale?
- Crash Handling /  
Supervision!

# Hmm...

## Scaling

- Just add More Processes

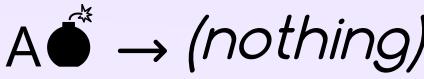
## Fault Tolerance

- **Keep track** of other processes
- **Restart things** when they fail

# Keeping Track: Links, Monitors

Crash Handling  
comes in two (and a  
*half*) kinds:

A Monitors B



A Linked to B



A Linked to B; A trapping exits



# Restart when it fails: Supervision

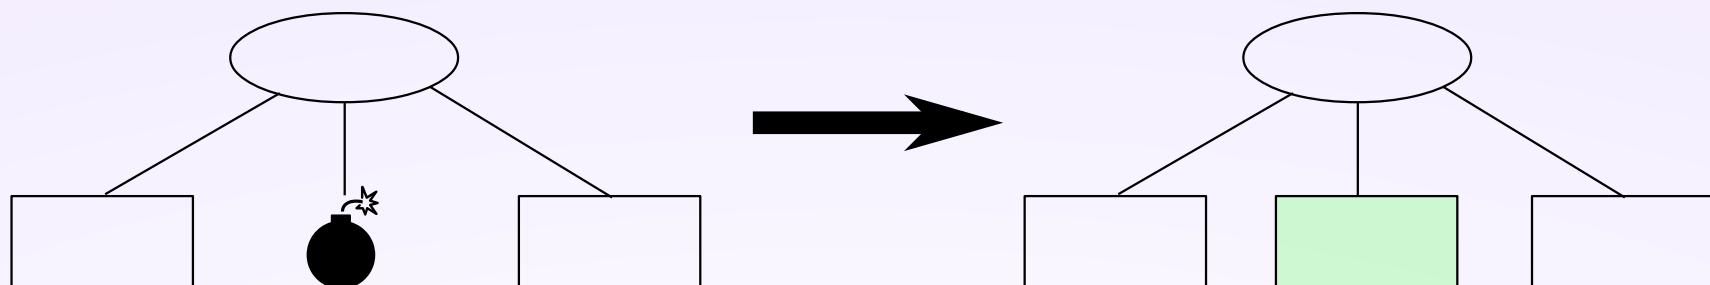
- **Workers** do work
- **Supervisors** supervise
  - Start workers
  - Restart workers *with known state* in case of a failure.
  - When too many failing workers in short time, terminate themselves.
- *Customizable!*
  - **Supervision strategies:** “What workers should be restarted when one fails”
  - How many failing workers in X seconds are considered ‘too many’?

# Supervision Strategies



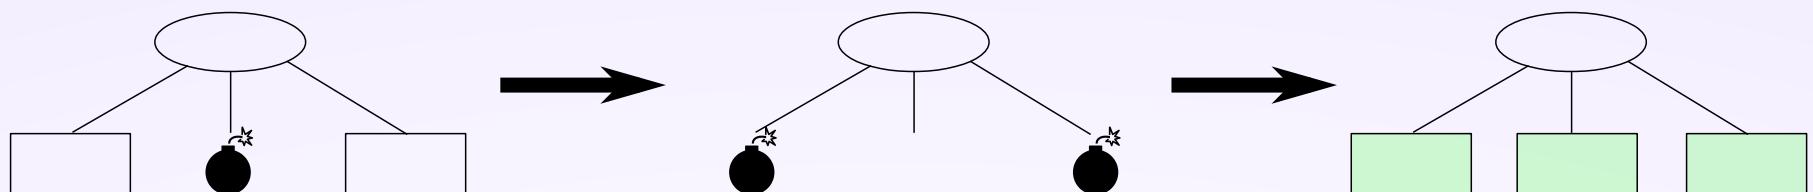
# Supervision Strategies: One For One

- When a worker dies, it is restarted.



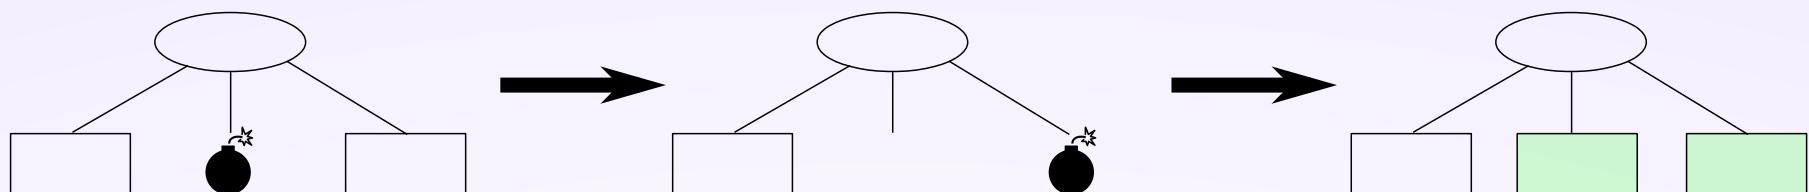
# Supervision Strategies: One For All

- When a worker dies, *all* of my workers are restarted.
- Useful if workers depend on each other.



# Supervision Strategies: Rest For One

- When a worker dies, that one *and all newer ones* are restarted.
- Useful if processes are in some kind of chain.



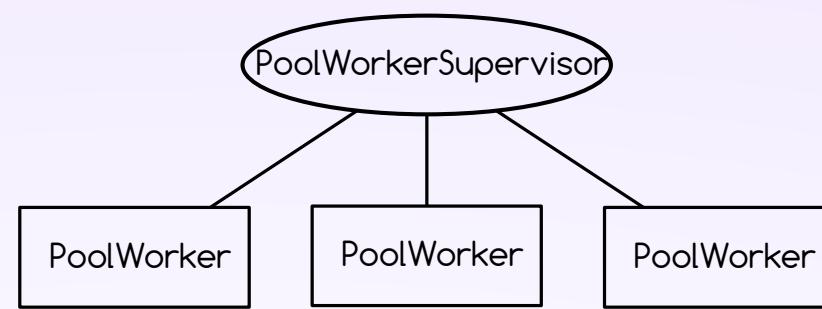
Supervision Strategies:  
**'Simple' One For One**

- *is not like the others*
- Can only use workers of single type
- Does not start out with list of workers, but with *blueprint*.
- Useful for dynamically changing how many workers.

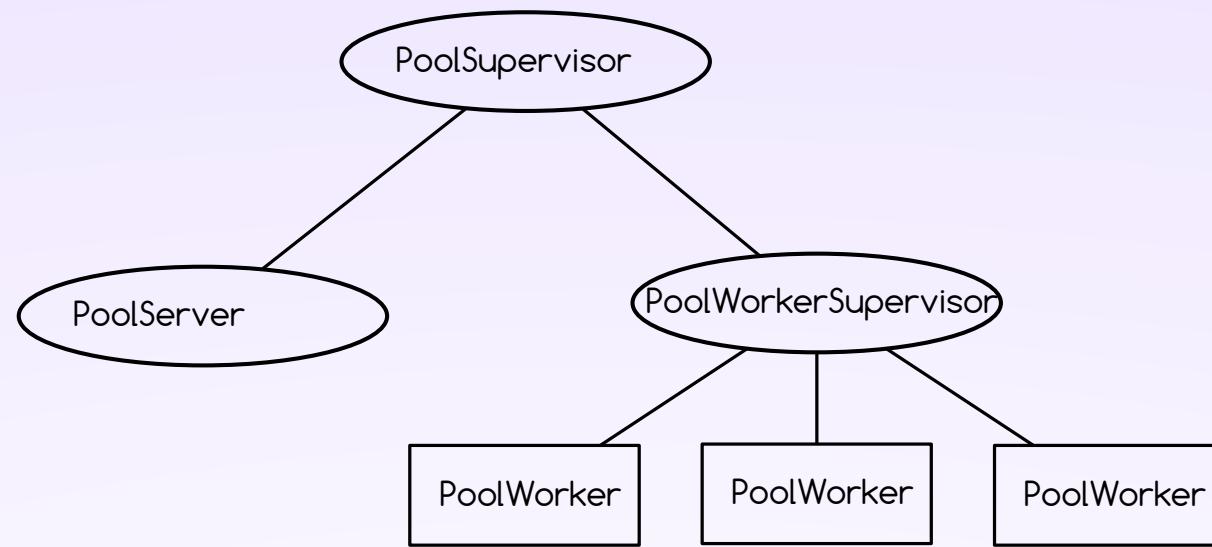
# Supervision Trees

- Layers of Fault Tolerance
- Proper cleanup (prevents floating processes/memory leaks).

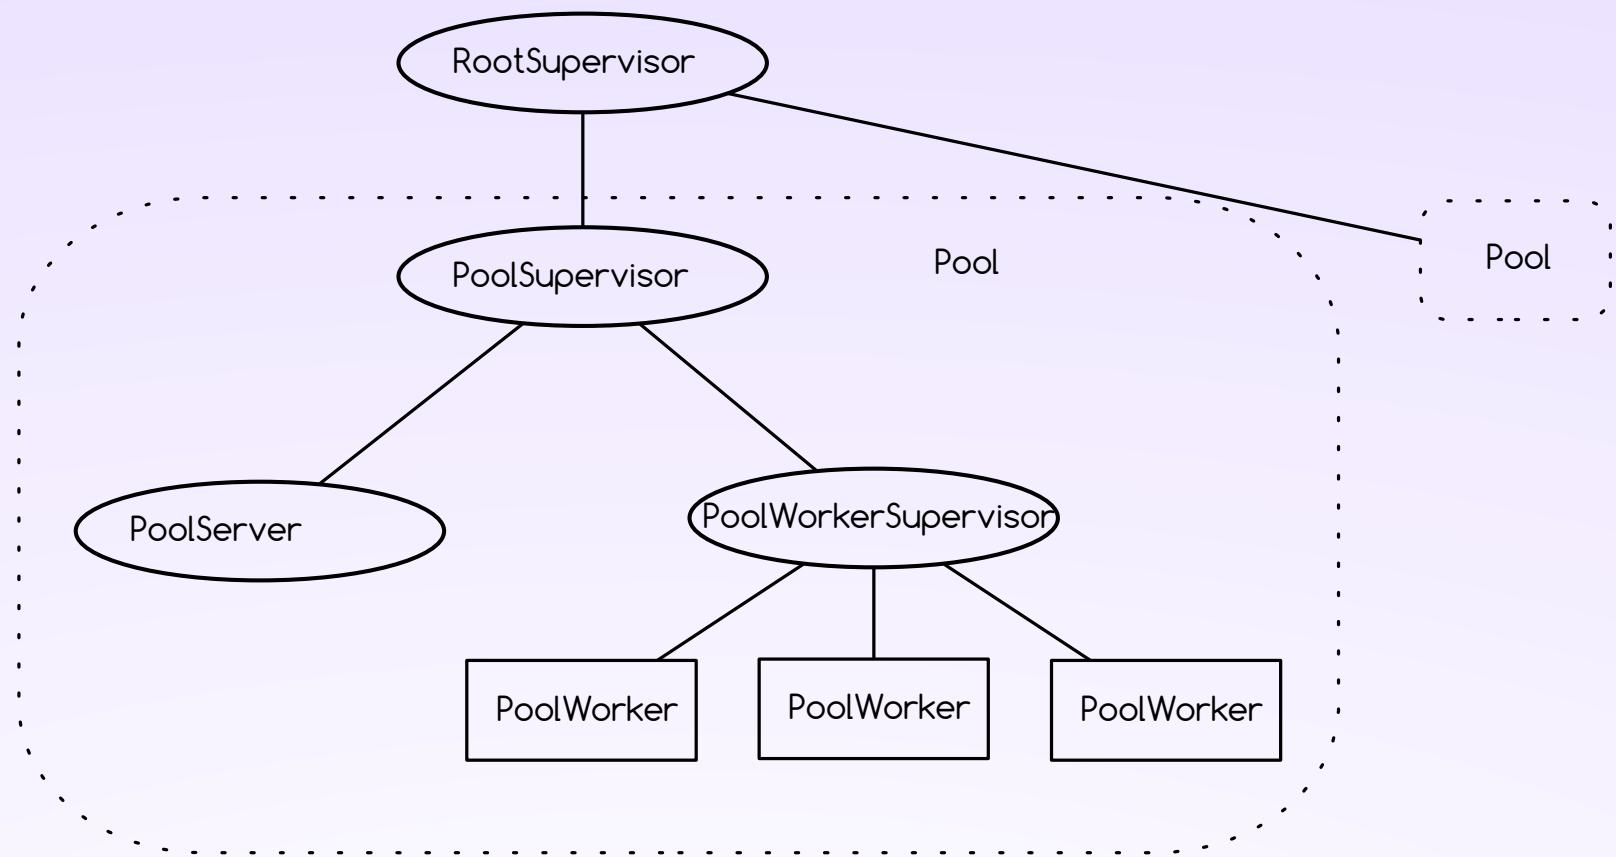
# Supervision Trees: Example



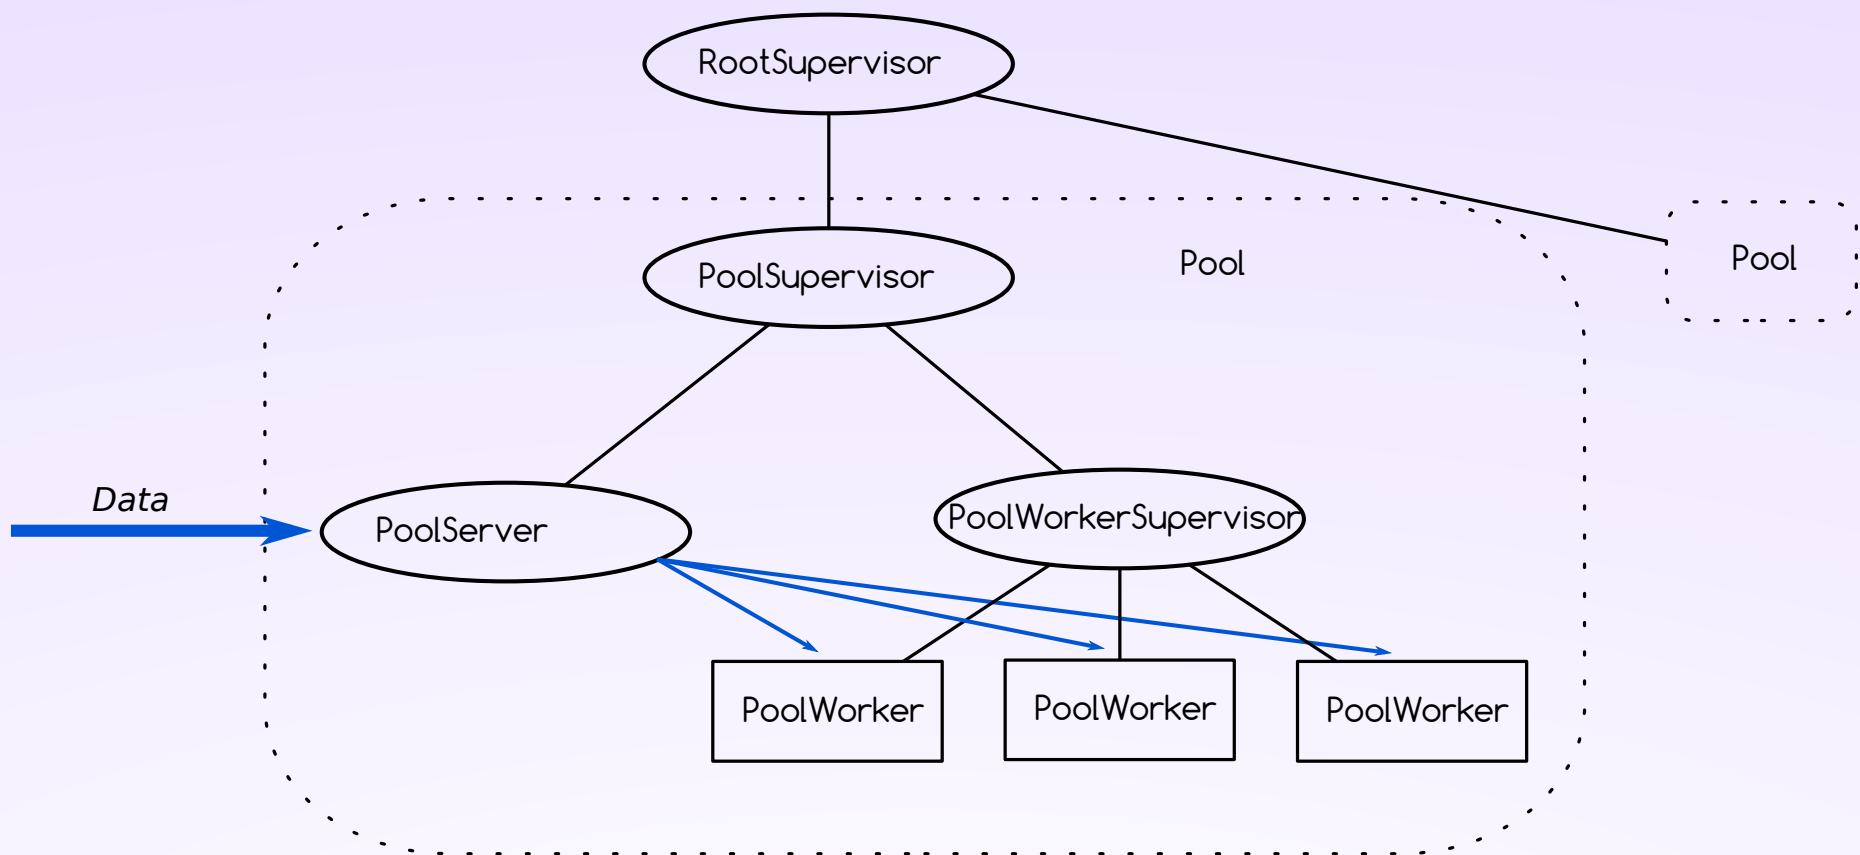
# Supervision Trees: Example



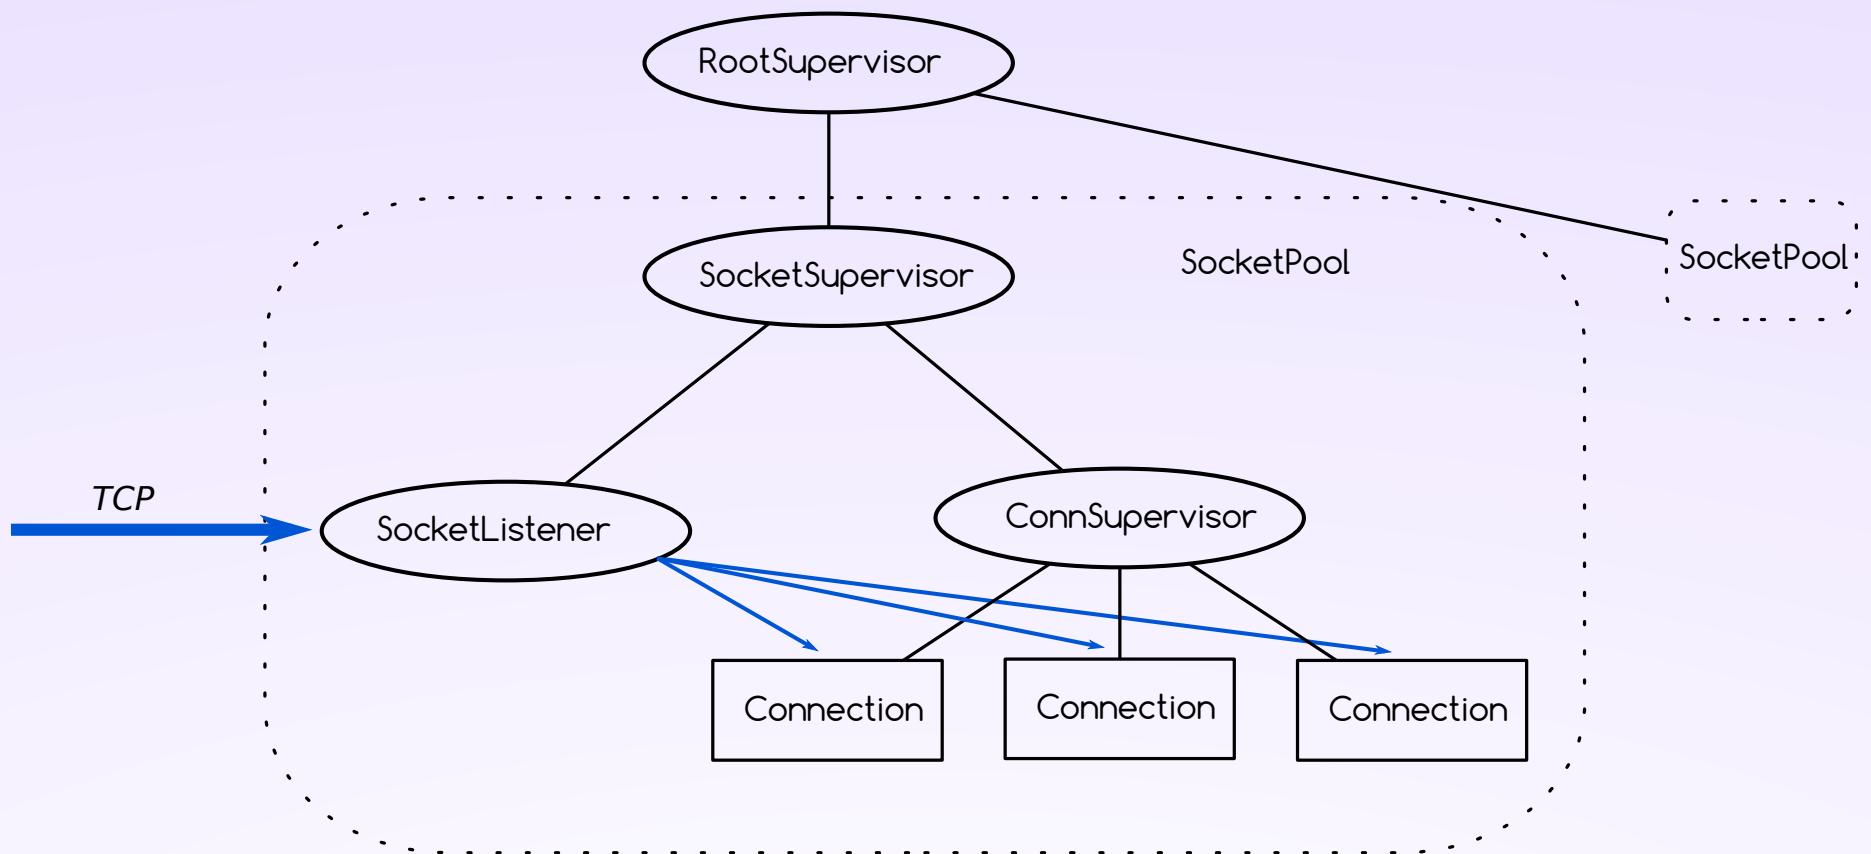
# Supervision Trees: Example



# Supervision Trees: Example



# Supervision Trees: Example

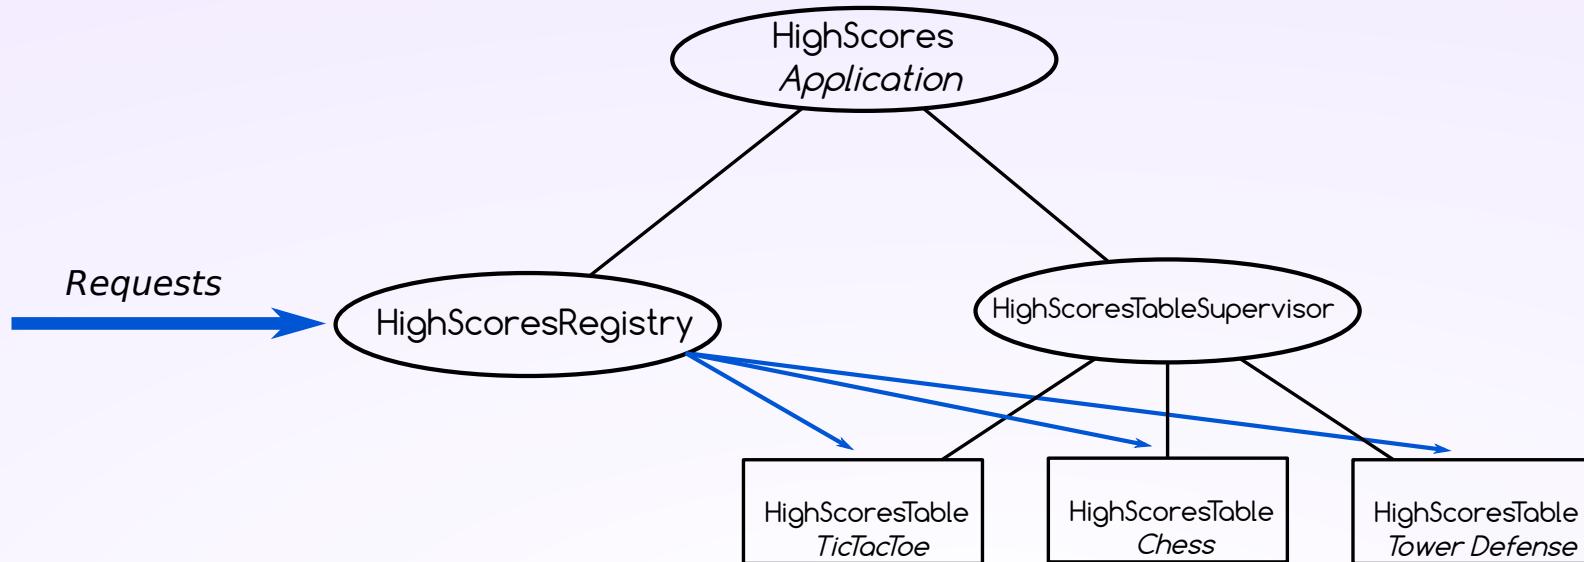


# Application

- Way to *package* a Supervision Tree
- *Stop/start/reload* as a unit.
- *Configurable* as a unit.
- *Re-usable* in other systems.
- Nice Debugging → :observer.start

# High Scores: Multiple games

- Add our HighScoresTable as worker in a Supervision Tree
- Add Server that maps game names to HighScoreTables.



```

defmodule HighScoresTable do
  use GenServer
  @moduledoc "A process that stores a table of high scores."
  @doc "A table is list of scores and its max. length"
  defstruct scores: [], max_players: nil

  defmodule Score do
    @doc "A single score"
    defstruct name: "", score: 0, time: nil
  end

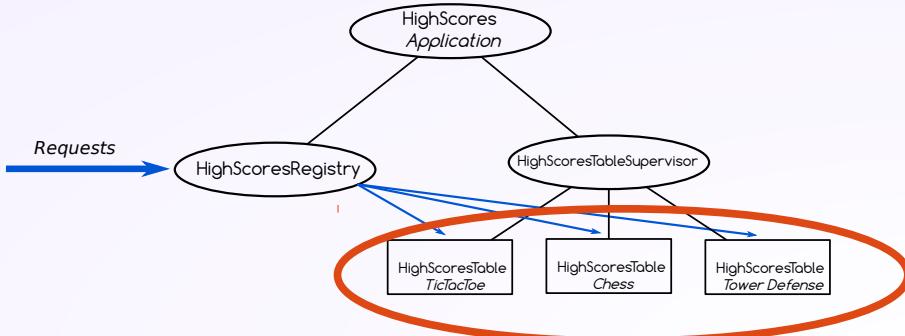
  # Outward API

  @doc "Starts the process, internally calls `init`."
  def start_link(max_players) do
    GenServer.start_link(__MODULE__, max_players)
  end

  @doc "Adds a new `score` for `name` to the table."
  def add_score(pid, name, score) do
    time = :erlang.time |> Time.from_erl!
    score = %Score{name: name, score: score, time: time}
    GenServer.call(pid, {:add_score, score})
  end

  @doc "Lists current scores"
  def current_scores(pid) do
    GenServer.call(pid, :current_scores)
  end

```



```

  # internal GenServer callbacks

  @doc "Initializes the GenServer"
  def init(max_players) do
    initial_state = %HighScoresTable{max_players: max_players}
    {:ok, initial_state}
  end

  @doc "Adds a score to the table, updating the internal state."
  def handle_call({:add_score, score = %Score{}}, _from, state) do
    {response, new_scores} =
    maybe_add_score(state.scores, state.max_players, score)
    new_state = %HighScoresTable{state | scores: new_scores}
    {:reply, response, new_state}
  end

  @doc "Lists the current scores"
  def handle_call(:current_scores, _from, state) do
    {:reply, state.scores, state}
  end

  defp maybe_add_score(scores, max_players, new_score) do
    split_scores = Enum.split_while(scores, fn score ->
      score.score >= new_score.score end)
    case split_scores do
      {all, []} when length(all) >= max_players ->
        {:unfortunate, scores}
      {higher, lower} ->
        new_scores = (higher ++ [new_score] ++ lower) |>
        Enum.take(max_players)
        {{:congrats, length(higher)+1}, new_scores}
    end
  end
end

```

```

defmodule HighScoresRegistry do
  use GenServer

  @doc "Starts the HighScoresRegistry, with the module name as
process alias so we don't have to pass the pid around"
  def start_link do
    GenServer.start_link(__MODULE__, :ok, name: __MODULE__)
  end

  def init(:ok) do
    {:ok, %{}}
  end

  # Outward API

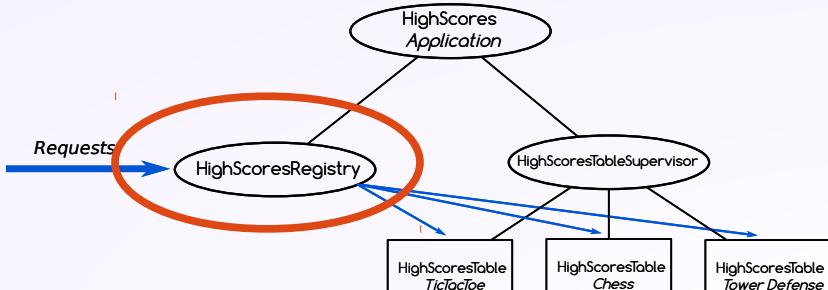
  @doc "Adds a new game"
  def add_game(game, max_players \\ 10) do
    GenServer.call(__MODULE__, {:add_game, game, max_players})
  end

  @doc "Adds `name`'s `score` on `game` to that games table"
  def add_score(game, name, score) do
    {:ok, hst_pid} = get_game_pid(game)
    HighScoresTable.add_score(hst_pid, name, score)
  end

  @doc "Lists the current scores for `game`"
  def current_scores(game) do
    {:ok, hst_pid} = get_game_pid(game)
    HighScoresTable.current_scores(hst_pid)
  end

  # Returns the pid of a certain game.
  defp get_game_pid(game) do
    GenServer.call(__MODULE__, {:get_game, game})
  end

```



```

# internal GenServer callbacks

  def handle_call({:add_game, game_name, max_players},
  _from, games) do
    if games[game_name] do
      {:reply, {:error, :game_already_exists}, games}
    else
      {:ok, pid} =
        HighScoresTableSupervisor.start_child([max_players])
      Process.monitor(pid)
      games = Map.put_new(games, game_name, {pid,
      max_players})
      {:reply, :ok, games}
    end
  end

  def handle_call({:get_game, game}, _from, games) do
    if games[game] do
      {pid, _} = games[game]
      {:reply, {:ok, pid}, games}
    else
      {:reply, {:error, :unknown_game}, games}
    end
  end

  @doc "Remove table processes when they crash, and start a
new table for that game."
  def handle_info({:DOWN, ref, :process, crashed_pid,
  _reason}, games) do
    games =
      for {game_name, {pid, max_players}} <- games, into: %
    {} do
      if pid == crashed_pid do
        {:ok, new_pid} =
          HighScoresTableSupervisor.start_child([max_players])
        {game_name, {new_pid, max_players}}
      else
        {game_name, {pid, max_players}}
      end
    end
    {:noreply, games}
  end
end

```

```

defmodule HighScoresTableSupervisor do
  use Supervisor

  def start_link do
    Supervisor.start_link(__MODULE__, :ok, name: __MODULE__)
  end

  def init(:ok) do
    children = [
      worker(HighScoresTable, [])
    ]

    supervise(children, strategy: :simple_one_for_one)
  end

  def start_child(args) do
    Supervisor.start_child(__MODULE__, args)
  end
end

```

```

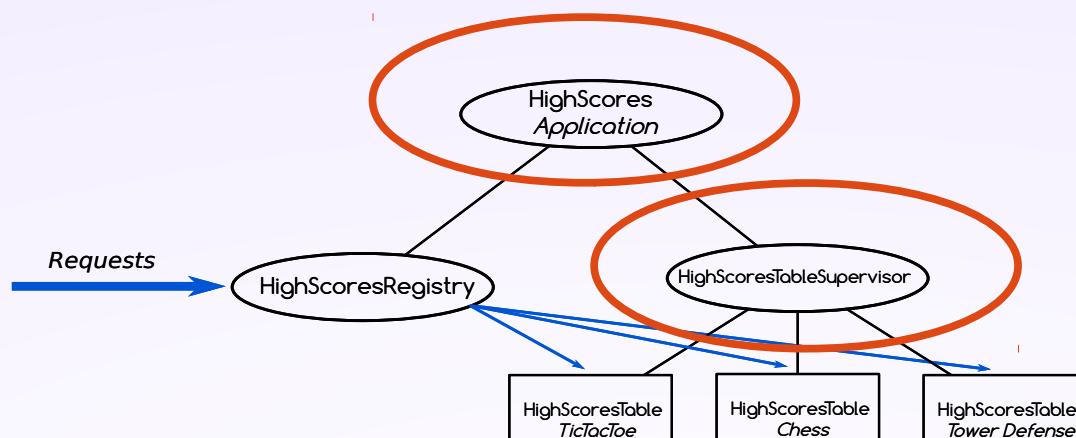
defmodule HighScores do
  use Application

  def start(_type, _args) do
    import Supervisor.Spec, warn: false

    # Define workers and child supervisors to be supervised
    children = [
      supervisor(HighScoresTableSupervisor, []),
      worker(HighScoresRegistry, [])
    ]

    opts = [strategy: :one_for_one, name: HighScores.Supervisor]
    Supervisor.start_link(children, opts)
  end
end

```



*-Live Demonstration-*

# Ways this could be improved

- Advanced Process naming using pg2.
- Storing things with ETS.
- Reduce more GenServer boilerplate using ExActor.
- Distribute score tables on multiple servers using GenServer.`multi_call`

# Conclusions

- It's possible to quickly build scalable, concurrent applications in Elixir.
- Elixir has many built-in tools and mechanisms to make this easy.
- There's a lot more you can do.

# Final Conclusion

- The most important thing, is to know what technologies are out there
- Use the right tool for the right job



# Further Learning

- Book: Elixir in Action – *Saša Jurić*  
<https://manning.com/books/elixir-in-action>
- Book: Learn You Some Erlang – Fred Hébert  
<http://learnyousomeerlang.com/>
- Internet: Elixir Documentation  
<http://elixir-lang.org/docs.html>
- Talk: Discovering Processes – *Saša Jurić*  
[https://youtu.be/y\\_b6RTes83c](https://youtu.be/y_b6RTes83c)

**CONGLATURATION !!!**  
**YOU HAVE COMPLETED**  
**A GREAT GAME.**  
**AND PROOVED THE JUSTICE**  
**OF OUR CULTURE.**  
**NOW GO AND REST OUR**  
**HEROES !**